Assignment 1 (ML for TS) - MVA

Firstname Lastname youremail1@mail.com
Firstname Lastname youremail2@mail.com

October 24, 2024

1 Introduction

Objective. This assignment has three parts: questions about convolutional dictionary learning, spectral features, and a data study using the DTW.

Warning and advice.

- Use code from the tutorials as well as from other sources. Do not code yourself well-known procedures (e.g., cross-validation or k-means); use an existing implementation.
- The associated notebook contains some hints and several helper functions.
- Be concise. Answers are not expected to be longer than a few sentences (omitting calculations).

Instructions.

- Fill in your names and emails at the top of the document.
- Hand in your report (one per pair of students) by Tuesday 28$^{\text{th}}$ October 23:59 PM.
- Rename your report and notebook as follows:
  `FirstnameLastname1_FirstnameLastname2.pdf` and
  `FirstnameLastname1_FirstnameLastname2.ipynb`.
  For instance, `LaurentOudre_CharlesTruong.pdf`.
- Upload your report (PDF file) and notebook (IPYNB file) using this link: LINK.

2 Convolution dictionary learning

Question 1

Consider the following Lasso regression:

$$\min_{\beta \in \mathbb{R}^p} \frac{1}{2} \|y - X\beta\|_2^2 \quad + \quad \lambda \|\beta\|_1 \tag{1}$$

where $y \in \mathbb{R}^n$ is the response vector, $X \in \mathbb{R}^{n \times p}$ the design matrix, $\beta \in \mathbb{R}^p$ the vector of regressors and $\lambda > 0$ the smoothing parameter.

Show that there exists $\lambda_{\max}$ such that the minimizer of (1) is $\mathbf{0}_p$ (a p-dimensional vector of zeros) for any $\lambda > \lambda_{\max}$.

Answer 1

$$\lambda_{\max} = \dots \tag{2}$$

Question 2

For a univariate signal $\mathbf{x} \in \mathbb{R}^n$ with n samples, the convolutional dictionary learning task amounts to solving the following optimization problem:

$$\min_{(\mathbf{d}_k)_k, (\mathbf{z}_k)_k \| \mathbf{d}_k \|_2^2 \leq 1} \quad \left\| \mathbf{x} - \sum_{k=1}^{K} \mathbf{z}_k * \mathbf{d}_k \right\|_2^2 \quad + \quad \lambda \sum_{k=1}^{K} \|\mathbf{z}_k\|_1 \tag{3}$$

where $\mathbf{d}_k \in \mathbb{R}^L$ are the K dictionary atoms (patterns), $\mathbf{z}_k \in \mathbb{R}^{N-L+1}$ are activations signals, and $\lambda > 0$ is the smoothing parameter.

Show that

- for a fixed dictionary, the sparse coding problem is a lasso regression (explicit the response vector and the design matrix);
- for a fixed dictionary, there exists $\lambda_{\max}$ (which depends on the dictionary) such that the sparse codes are only 0 for any $\lambda > \lambda_{\max}$.

Answer 2

On sait que :

$$(z_k * d_k)_i = \sum_{j=1}^{L} z_{k,j} d_{k,j-i}$$

$$= z_{k,1} d_{k,i} + z_{k,2} d_{k,i-1} + \dots + z_{k,L} d_{k,i-L+1}$$

Cela peut s'écrire comme un produit entre la matrice :

$$D_k = \begin{pmatrix} d_{k,1} & 0 & 0 & \cdots & 0 \\ d_{k,2} & d_{k,1} & 0 & \cdots & 0 \\ d_{k,3} & d_{k,2} & d_{k,1} & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ d_{k,L} & d_{k,L-1} & d_{k,L-2} & \cdots & 0 \\ 0 & d_{k,L} & d_{k,L-1} & \cdots & 0 \\ 0 & 0 & d_{k,L} & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & d_{k,L} \end{pmatrix}$$

et le vecteur $z_k \in \mathbb{R}^{N-L+1}$.

D_k est de dimensions $N \times (N - L + 1)$.

Comme on a une somme de convolution, on définit alors la matrice formée par les concaténations des D_k :

$$D = (D_1 D_2 D_3 \dots D_K)$$

et le vecteur $z = (z_1^T, z_2^T, \dots, z_K^T)^T$.

On a donc :

$$(3) \quad \Leftrightarrow \quad \min_{z \in \mathbb{R}^K} \frac{1}{2}\|X - Dz\|_2^2 + \lambda \|f\|_1$$

D'après la question précédente et sa réponse, on sait que :

A FINIR

3 Spectral feature

Let X_n $(n = 0, \dots, N-1)$ be a weakly stationary random process with zero mean and autocovariance function $\gamma(\tau) := \mathbb{E}(X_n X_{n+\tau})$. Assume the autocovariances are absolutely summable, i.e. $\sum_{\tau \in \mathbb{Z}} |\gamma(\tau)| < \infty$, and square summable, i.e. $\sum_{\tau \in \mathbb{Z}} \gamma^2(\tau) < \infty$. Denote the sampling frequency by f_s, meaning that the index n corresponds to the time n/f_s. For simplicity, let N be even.

The *power spectrum* S of the stationary random process X is defined as the Fourier transform of the autocovariance function:

$$S(f) := \sum_{\tau=-\infty}^{+\infty} \gamma(\tau) e^{-2\pi f \tau / f_s}. \tag{4}$$

The power spectrum describes the distribution of power in the frequency space. Intuitively, large values of $S(f)$ indicate that the signal contains a sine wave at the frequency f. There are many estimation procedures to determine this important quantity, which can then be used in a machine-learning pipeline. In the following, we discuss the large sample properties of simple estimation procedures and the relationship between the power spectrum and the autocorrelation.

(Hint: use the many results on quadratic forms of Gaussian random variables to limit the number of calculations.)

Question 3

In this question, let X_n $(n = 0, \dots, N-1)$ be a Gaussian white noise.

- Calculate the associated autocovariance function and power spectrum. (By analogy with the light, this process is called "white" because of the particular form of its power spectrum.)

Answer 3

Comme X_n ($n = 0, \dots, N-1$) est un bruit blanc gaussien, alors

$\forall n \in [\![0, N-1]\!], \quad X_i \sim \mathcal{N}(0, \sigma^2) \quad$ et les X_i sont identiques et indépendants. Le signal est donc stationnaire.

On a donc :

$$\gamma(\tau) = \mathbb{E}(X_n X_{n+\tau}) = \begin{cases} \mathbb{E}(X_n)\mathbb{E}(X_{n+\tau}) = 0 & \text{si } \tau \neq 0 \quad \text{(car les v.a. sont indép.)} \\ \mathbb{E}(X_n^2) = \mathrm{Var}(X_n) + (\mathbb{E}(X_n))^2 = \sigma^2 + 0 = \sigma^2 & \text{si } \tau = 0 \end{cases}$$

Donc :

$$\gamma(\tau) = \sigma^2 \cdot \mathbb{1}_{\{\tau = 0\}}$$

On calcule :

$$
\begin{aligned}
S(f) &= \sum_{\tau=-\infty}^{+\infty} \gamma(\tau) e^{-j2\pi f\tau/f_s} \\
&= \sigma^2 e^{-2\pi f\times 0/f_s} = \sigma^2
\end{aligned}
$$

Question 4

A natural estimator for the autocorrelation function is the sample autocovariance

$$
\hat{\gamma}(\tau) := (1/N) \sum_{n=0}^{N-\tau-1} X_n X_{n+\tau} \tag{5}
$$

for $\tau = 0, 1, \ldots, N-1$ and $\hat{\gamma}(\tau) := \hat{\gamma}(-\tau)$ for $\tau = -(N-1), \ldots, -1$.

- Show that $\hat{\gamma}(\tau)$ is a biased estimator of $\gamma(\tau)$ but asymptotically unbiased. What would be a simple way to de-bias this estimator?

Answer 4

On calcule l'espérance :

$$
\begin{aligned}
\mathbb{E}\left(\hat{\gamma}(\tau)\right) &= \mathbb{E}\left(\frac{1}{N}\sum_{n=0}^{N-\tau-1} X_n X_{n+\tau}\right) \\
&= \frac{1}{N}\sum_{n=0}^{N-\tau-1} \mathbb{E}(X_n X_{n+\tau}) \\
&= \frac{1}{N}\sum_{n=0}^{N-\tau-1} \gamma(\tau)
\end{aligned}
$$

Comme le processus est stationnaire, $\mathbb{E}(X_n X_{n+\tau})$ (et donc $\gamma(\tau)$) ne dépendent pas de n (seul l'écart τ importe), et donc :

$$
\begin{aligned}
\mathbb{E}\left(\hat{\gamma}(\tau)\right) &= \frac{1}{N}\gamma(\tau)\sum_{n=0}^{N-\tau-1} 1 \\
&= \frac{N-\tau}{N}\gamma(\tau)
\end{aligned}
$$

On n'a pas $\mathbb{E}\left(\hat{\gamma}(\tau)\right) = \gamma(\tau)$, donc l'estimateur est biaisé.

En revanche :

$$
\frac{N-\tau}{N} \xrightarrow[N\to\infty]{} 1
$$

Donc :

$$
\lim_{N\to\infty} \mathbb{E}\left(\hat{\gamma}(\tau)\right) = \gamma(\tau)
$$

Donc l'estimateur est asymptotiquement sans biais.

Pour débiaiser l'estimateur, il suffit de le multiplier par $\frac{N}{N-\tau}$.

Question 5

Define the discrete Fourier transform of the random process $\{X_n\}_n$ by

$$J(f) := (1/\sqrt{N}) \sum_{n=0}^{N-1} X_n e^{-2\pi i f n / f_s} \tag{6}$$

The *periodogram* is the collection of values $|J(f_0)|^2, |J(f_1)|^2, \dots, |J(f_{N/2})|^2$ where $f_k = f_s k/N$. (They can be efficiently computed using the Fast Fourier Transform.)

- Write $|J(f_k)|^2$ as a function of the sample autocovariances.
- For a frequency f, define $f^{(N)}$ the closest Fourier frequency f_k to f. Show that $|J(f^{(N)})|^2$ is an asymptotically unbiased estimator of $S(f)$ for $f > 0$.

Answer 5

Question 6

In this question, let X_n ($n = 0, \dots, N-1$) be a Gaussian white noise with variance $\sigma^2 = 1$ and set the sampling frequency to $f_s = 1$ Hz

- For $N \in \{200, 500, 1000\}$, compute the *sample autocovariances* ($\hat{\gamma}(\tau)$ vs τ) for 100 simulations of X. Plot the average value as well as the average $\pm$, the standard deviation. What do you observe?
- For $N \in \{200, 500, 1000\}$, compute the *periodogram* ($|J(f_k)|^2$ vs f_k) for 100 simulations of X. Plot the average value as well as the average $\pm$, the standard deviation. What do you observe?

Add your plots to Figure 1.

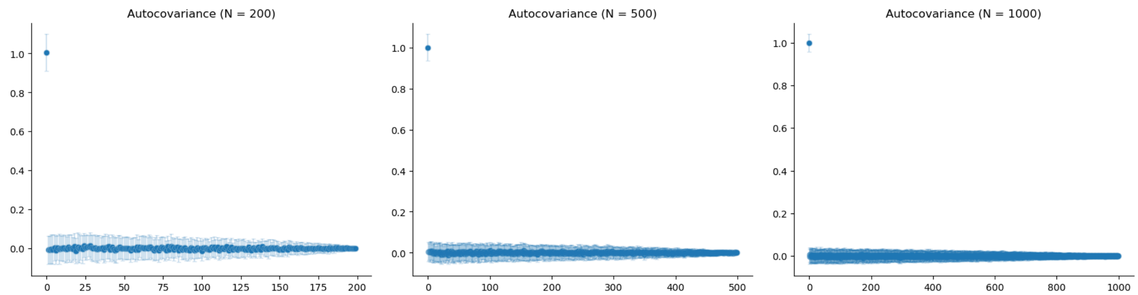

Autocovariance ($N = 200$) Autocovariance ($N = 500$) Autocovariance ($N = 1000$)

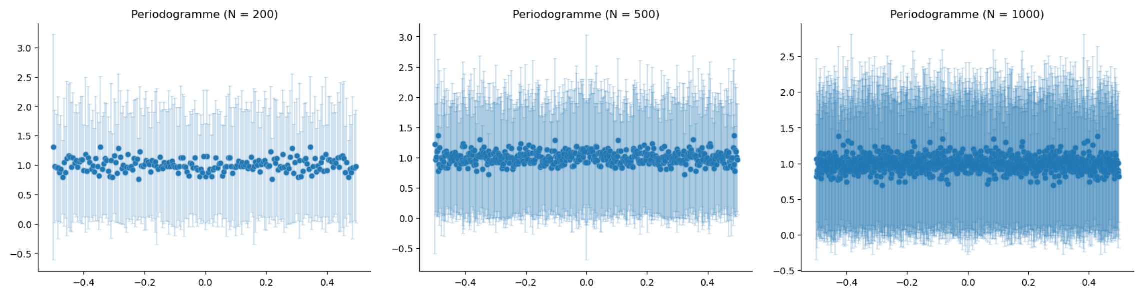

Periodogram ($N = 200$) Periodogram ($N = 500$) Periodogram ($N = 1000$)

Figure 1: Autocovariances and periodograms of a Gaussian white noise (see Question 6).

Answer 6

- Pour l'autocovariance, on remarque qu'on a à chaque fois une autocovariance égale à 0 pour $\tau = 0$, ce qui est logique puisqu'on a montré que : $\gamma(\tau) = \sigma^2 \cdot \mathbb{1}_{\{\tau=0\}}$ et ici $\sigma^2 = 1$.

  On remarque aussi que la valeur d'autocovariance d'échantillon est ensuite très proche de 0, ce qui est logique puisque on est en présence de bruit blanc gaussien et deux points successifs du signal sont donc indépendants et donc $\gamma(\tau) = 0$.

- Pour le spectrogramme, on remarque que toutes les fréquences sont autant représentées les unes que les autres, ce qui est logique puisqu'on est dans le cas de bruit blanc gaussien. On a représenté les fréquences négatives mais on aurait pu se limiter aux fréquences positives puisque la transformée de Fourier est symétrique pour les signaux réels. On a par ailleurs bien représenté des fréquences qui vont jusquà 0.5 Hz (soit jusqu'à la fréquence de Nyquist $f_s/2$. Comme les fréquences sont regardées sur l'ensemble du signal, on a toujours un écart-type de la densité spectrale de puissance identique pour chaque fréquence et ce peut importe la taille de N.

Question 7

We want to show that the estimator $\hat{\gamma}(\tau)$ is consistent, i.e. it converges in probability when the number N of samples grows to ∞ to the true value $\gamma(\tau)$. In this question, assume that X is a wide-sense stationary *Gaussian* process.

- Show that for $\tau > 0$

$$\mathrm{var}(\hat{\gamma}(\tau)) = (1/N) \sum_{n=-(N-\tau-1)}^{n=N-\tau-1} \left(1 - \frac{\tau + |n|}{N}\right) \left[\gamma^2(n) + \gamma(n-\tau)\gamma(n+\tau)\right]. \tag{7}$$

(Hint: if $\{Y_1, Y_2, Y_3, Y_4\}$ are four centered jointly Gaussian variables, then $\mathbb{E}[Y_1Y_2Y_3Y_4] = \mathbb{E}[Y_1Y_2]\mathbb{E}[Y_3Y_4] + \mathbb{E}[Y_1Y_3]\mathbb{E}[Y_2Y_4] + \mathbb{E}[Y_1Y_4]\mathbb{E}[Y_2Y_3]$.)

- Conclude that $\hat{\gamma}(\tau)$ is consistent.

Answer 7

On sait que :

$$\mathrm{Var}(\hat{\gamma}(\tau)) = \mathbb{E}(\hat{\gamma}(\tau)^2) - (\mathbb{E}(\hat{\gamma}(\tau)))^2$$

Donc on va calculer :

$$\hat{\gamma}(\tau)^2 = \frac{1}{N^2} \left(\sum_{n=0}^{N-\tau-1} X_n X_{n+\tau} \right)^2$$

$$= \frac{1}{N^2} \sum_{n=0}^{N-\tau-1} \sum_{m=0}^{N-\tau-1} X_n X_{n+\tau} X_m X_{m+\tau}$$

On passe à l'espérance :

$$\mathbb{E}\left(\hat{\gamma}(\tau)^2\right) = \mathbb{E}\left(\frac{1}{N^2} \sum_{n=0}^{N-\tau-1} \sum_{m=0}^{N-\tau-1} X_n X_{n+\tau} X_m X_{m+\tau} \right)$$

$$= \frac{1}{N^2} \sum_{n=0}^{N-\tau-1} \sum_{m=0}^{N-\tau-1} \mathbb{E}(X_n X_{n+\tau} X_m X_{m+\tau})$$

Comme $\{X_n, X_{n+\tau}, X_m, X_{m+\tau}\}$ sont 4 variables gaussiennes centrées conjointement et en utilisant l'indice donné dans la question, on a :

$$\begin{aligned}\mathbb{E}(X_n X_{n+\tau} X_m X_{m+\tau}) = \mathbb{E}(X_n X_{n+\tau})\mathbb{E}(X_m X_{m+\tau}) + \mathbb{E}(X_n X_m)\mathbb{E}(X_{n+\tau} X_{m+\tau}) \\ + \mathbb{E}(X_n X_{m+\tau})\mathbb{E}(X_{n+\tau} X_m)\end{aligned}$$

Comme $\gamma(\tau) = \mathbb{E}(X_m X_{m+\tau})$ et que $\gamma(\tau)$ ne dépend que de la distance temporelle τ entre la 1ère et la 2ème variable, alors on a :

$$\mathbb{E}(X_m X_{n+\tau} X_n X_{m+\tau}) = \gamma(\tau)\gamma(\tau) + \gamma(m-n)\gamma(m-n) + \gamma(m+\tau-n)\gamma(m-n-\tau)$$

Donc :

$$\mathbb{E}(\hat{\gamma}(\tau)^2) = \frac{1}{N^2} \sum_{n=0}^{N-\tau-1} \sum_{m=0}^{N-\tau-1} \left(\gamma^2(\tau) + \gamma^2(m-n) + \gamma(m+\tau-n)\gamma(m-n-\tau)\right)$$

$$= \frac{1}{N^2} \left(\gamma^2(\tau) \sum_{n=0}^{N-\tau-1} \sum_{m=0}^{N-\tau-1} 1 + \sum_{n=0}^{N-\tau-1} \sum_{m=0}^{N-\tau-1} \gamma^2(m-n) + \sum_{n=0}^{N-\tau-1} \sum_{m=0}^{N-\tau-1} \gamma(m+\tau-n)\gamma(m-n-\tau) \right)$$

A FINIR

Contrary to the correlogram, the periodogram is not consistent. It is one of the most well-known estimators that is asymptotically unbiased but not consistent. In the following question, this is proven for Gaussian white noise, but this holds for more general stationary processes.

Question 8

Assume that X is a Gaussian white noise (variance σ^2) and let $A(f) := \sum_{n=0}^{N-1} X_n \cos(-2\pi f n/f_s$ and $B(f) := \sum_{n=0}^{N-1} X_n \sin(-2\pi f n/f_s$. Observe that $J(f) = (1/N)(A(f) + iB(f))$.

- Derive the mean and variance of $A(f)$ and $B(f)$ for $f = f_0, f_1, \ldots, f_{N/2}$ where $f_k = f_s k/N$.
- What is the distribution of the periodogram values $|J(f_0)|^2$, $|J(f_1)|^2$, $\ldots$, $|J(f_{N/2})|^2$.
- What is the variance of the $|J(f_k)|^2$? Conclude that the periodogram is not consistent.
- Explain the erratic behavior of the periodogram in Question 6 by looking at the covariance between the $|J(f_k)|^2$.

Answer 8

Question 9

As seen in the previous question, the problem with the periodogram is the fact that its variance does not decrease with the sample size. A simple procedure to obtain a consistent estimate is to divide the signal into K sections of equal durations, compute a periodogram on each section, and average them. Provided the sections are independent, this has the effect of dividing the variance by K. This procedure is known as Bartlett's procedure.

- Rerun the experiment of Question 6, but replace the periodogram by Barlett's estimate (set $K = 5$). What do you observe?

Add your plots to Figure 2.

Answer 9

Avec la méthode de Bartlett, on voit que la variance est plus petite qu'avec la méthode du périodogramme classique. Il a cependant moins de points lorsque qu'on compare à N identique, ce qui est logique puisqu'on a travaillé sur des segments de taille plus petits. On ne voit pas de réelle diminution de l'écart type avec la taille de l'échantillon. C'est peut être parce ???????

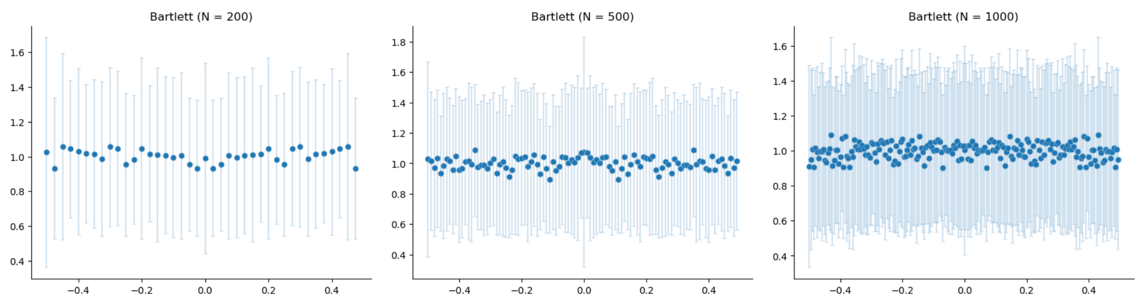

Periodogram ($N = 200$) Periodogram ($N = 500$) Periodogram ($N = 1000$)

Figure 2: Barlett's periodograms of a Gaussian white noise (see Question 9).

4 Data study

4.1 General information

Context. The study of human gait is a central problem in medical research with far-reaching consequences in the public health domain. This complex mechanism can be altered by a wide range of pathologies (such as Parkinson's disease, arthritis, stroke,...), often resulting in a significant loss of autonomy and an increased risk of falls. Understanding the influence of such medical disorders on a subject's gait would greatly facilitate early detection and prevention of those possibly harmful situations. To address these issues, clinical and bio-mechanical researchers have worked to objectively quantify gait characteristics.

Among the gait features that have proved their relevance in a medical context, several are linked to the notion of step (step duration, variation in step length, etc.), which can be seen as the core atom of the locomotion process. Many algorithms have, therefore, been developed to automatically (or semi-automatically) detect gait events (such as heel-strikes, heel-off, etc.) from accelerometer and gyrometer signals.

Data. Data are described in the associated notebook.

4.2 Step classification with the dynamic time warping (DTW) distance

Task. The objective is to classify footsteps and then walk signals between healthy and non-healthy.

Performance metric. The performance of this binary classification task is measured by the F-score.

Question 10

Combine the DTW and a k-neighbors classifier to classify each step. Find the optimal number of neighbors with 5-fold cross-validation and report the optimal number of neighbors and the associated F-score. Comment briefly.

Answer 10

Question 11

Display on Figure 3 a badly classified step from each class (healthy/non-healthy).

Answer 11

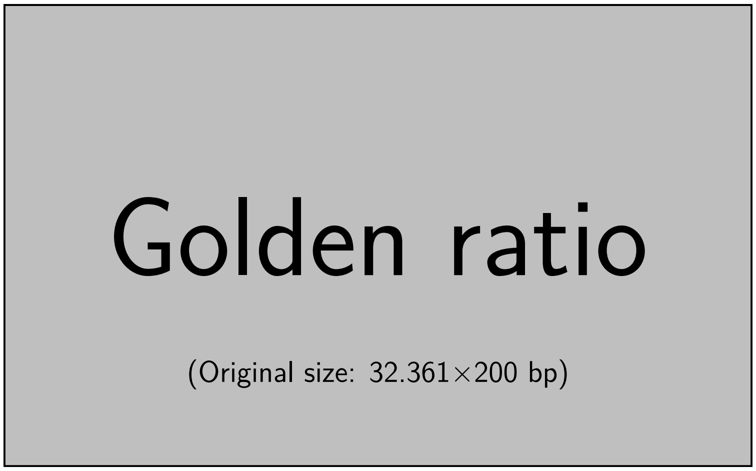

Badly classified healthy step

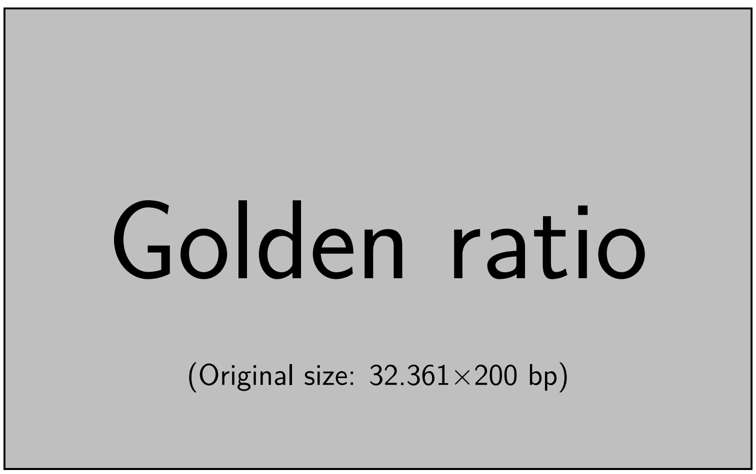

Badly classified non-healthy step

Figure 3: Examples of badly classified steps (see Question 11).